NLP Homework

Data and code given to you:

- `glove.6B.50d.10k.mat`: A set of embeddings from GLoVE of 10,000 words.
- `word-sim.csv`: Human rated scores of word similarity.
- `readWordVectors.m`: A function to extract words and embeddings.
- `word2vec.m`: A function to look up the embedding of a single word.
- `testAnalogies.m`: A function to test your `analogy.m` code.

Usage:

```
[vectors, words] = readWordVectors('glove.6B.50d.10k.mat');
embedding = word2vec(vectors, words, 'shoe');
```

1 Word Similarity

Implement the function `wordSimilarity.m` which computes the cosine similarity between two given words.

Test it on a few words of your choice to make sure that the results are reasonable.

For example, "cat" and "pet" should return a high value (close to 1.0), while "jump" and "coast" should return a lower value.

Now, implement the function `computeBehaviorCorrelation.m`. In this function, you will compare human rated similarities with the word embedding similarity measure you implemented in `wordSimilarity`. Follow the comments in the code to return the Spearman correlation. You should get a value of 0.6298.

2 Closest Words

Implement the function `mostSimilar.m` that will return the 10 closest words (in terms of cosine distance), to the given embedding.

Hint: The `norm` and `vecnorm` functions may be useful.

When that is working, implement `plotMostSimilar.m` to plot a word and its 30 closest words. You will need to project the embeddings onto 2D space using

PCA.

3 Analogies

Finally, implement the `analogy.m`. Use `testAnalogies.m` to run your function through a set of analogies. You should get 13 out of 13 correct.

Next, find 2 analogies that don't work and 2 additional ones that work (they must be different than those given in `testAnalogies.m`.

4 Clustering

This is an exploratory analysis. Run KMeans on the vectors with different number of clusters. Write a paragraph or two on the clusters you found, whether they're meaningful of not.

5 Reference/Acknowledgements

This is just for reference.

- GLoVE embeddings: https://nlp.stanford.edu/projects/glove/
- Word similarity: http://alfonseca.org/eng/research/wordsim353.html